

# Report Topics for Computational Linguistics & NLP

-Liviu P Dinu & Ana Uban-

Topics or projects:

1. Participate in shared tasks and competitions in the field of NLP (Kaggle is not accepted - if you need datasets start [here](#)): [SemEval](#), [CLEF](#), [PAN](#), [VarDial](#), any shared tasks associated with [top ranking](#) (A and A\* according to [core](#)) NLP conferences (EMNLP, COLING, ACL, NAACL, EACL, LREC, etc.)
2. Automatic text summarization (abstractive or extractive)
  - Ideas: a news summarizer/generating news headlines (option: for Romanian - more challenging for abstractive), summarizer for scientific articles/generating the abstracts automatically), summarizer for stories/literature; + evaluate it using standard metrics and report results
  - Papers: [Abstractive Summarization: A Survey of the State of the Art](#) (2019), [Recent automatic text summarization techniques: a survey](#) (2016)
3. POS-tagging (part of speech tagging)
  - Ideas: implement a POS-tagging algorithm from scratch (optional: include graphical visualization) + evaluate and report results
  - Papers: [POS Tagging for Arabic Tweets](#), [Non-lexical neural architecture for fine-grained POS Tagging](#)
  - Existing tools: [The Stanford POS Tagger](#)
  - Annotated datasets: Penn Treebank from NLTK
4. Named Entity Recognition and other Information Extraction tasks
  - Ideas: implement a NER algorithm from scratch, medical NER
  - Papers: [A survey of named entity recognition and classification](#), [Results of the WNUT2017 Shared Task on Novel and Emerging Entity Recognition](#), [Enhancing clinical concept extraction with contextual embeddings](#), [OntoNERdIE – Mapping and Linking Ontologies to Named Entity Recognition and Information Extraction Resources](#)
  - Applications: [Lark](#) needed it for food parsing, Bing needs it for search, processing medical text, populating knowledge bases
  - Existing tools: implemented in spacy, [YODIE Named Entity Disambiguation \(English\)](#)
  - Annotated datasets: [juand-r/entity-recognition-datasets](#), Tweets [here](#)
5. Syntax Parsing; Text processing at the syntax level (dependency grammars/dependency parsing; constituent parsing)
  - Ideas: implement a syntax parser from scratch (optional: on Romanian/rare language), create a graphical visualization of parsed sentence
  - Papers: [Accurate Unlexicalized Parsing](#), [Proceedings of the...](#), [Parsing as Sequence Labeling](#)
  - Existing tools: [Stanford parser](#)
  - Applications: [Grammarly](#), [Grammarly Knock-off](#)

- Annotated datasets: Penn Treebank in NLTK
- 6. Corpus/dataset creation (collection, cleaning, annotation, etc; e.g. [Twitter/Reddit API](#), [web scraper for news articles/political speeches/meeting transcripts/dictionaries](#), ...) (look at [LREC](#), benchmark dataset type papers) - either create new kind of corpus OR complement corpus collection with some form of linguistic analysis
  - Ideas: corpus of Romanian product/movie reviews, annotate with sentiment; corpus of news, annotate with emotions expressed; corpus of non-English Tweets, annotate with optimism/pessimism/mental health ("I am diagnosed with depression");
- 7. Text simplification
  - Ideas: implement a text simplification solution + evaluate and report results
  - Applications: Simple Wikipedia, language learning
  - Papers: [Exploring neural text simplification methods](#)
  - Annotated datasets: Simple Wikipedia,
- 8. Textual semantic similarity, text clustering; NLU (natural language inference, entailment)
  - Ideas: implement a textual entailment model + evaluate; train document/sentence embeddings for semantic similarity;
  - Papers: [Distributed Representations of Sentences and Documents \(doc2vec\)](#), [Recognizing Textual Entailment in Twitter Using Word Embeddings](#), [Siamese recurrent architectures for learning sentence similarity](#), [Word n-gram attention models for sentence similarity and inference](#)
  - Annotated datasets: [SNLI](#) (for entailment);
- 9. Distributional semantics, word embeddings, contextual embeddings
  - Ideas: implement, compare and evaluate various measures of similarity metrics on embeddings; visualization tool for embedding spaces; train embeddings on new domain and evaluate/discuss - needs large dataset (e.g. embeddings for medical data, embeddings for social media slang); evaluate and compare methods for embeddings compositionality
  - Papers: [Distributed Representations of Words and Phrases and their Compositionality](#), [Enriching Word Vectors with Subword Information](#), [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#), [\[1802.04302\] Evaluating Compositionality in Sentence Embeddings](#), [How Contextual are Contextualized Word Representations? Comparing the Geometry of BERT, ELMo, and GPT-2 Embeddings](#)
  - Existing tools: [Embedding projector - visualization of high-dimensional data](#), [GloVe: Global Vectors for Word Representation](#) (pre-trained GloVe embeddings), <https://github.com/facebookresearch/MUSE> (multilingual FastText embeddings trained on Wikipedia)
  - Datasets of annotated word similarity: [SimLex](#), [WordSim](#)
- 10. Fake news detection, rumor detection, propaganda detection
  - Examples: implement a fake news detection system + evaluate; implement automatic fact-checker (e.g. like <https://www.factual.ro/>); implement clickbait detector; detector of fake news in specific domain (political, medical)
  - Applications: [pheme](#), Twitter/Instagram/FB integrated fake news detection

- Papers: [“Liar, Liar Pants on Fire”: A New Benchmark Dataset for Fake News Detection](#), [We Built a Fake News / Click Bait Filter: What Happened Next Will Blow Your Mind!](#), [SemEval-2019 Task 7: RumourEval. Determining Rumour Veracity and Support for Rumours](#), [SemEval-2020 Task 11: Detection of Propaganda Techniques in News Articles](#), [Can We Spot the "Fake News" Before It Was Even Written?](#),
11. Deception detection, style transfer
- Ideas: system that automatically detects impersonation attempts in published long texts; generator of text in the style of a given author (challenging); identify impersonators on social media: discrepancies between declared age/sex and real one
  - Papers: [Pastiche Detection Based on Stopword Rankings. Exposing Impersonators of a Romanian Writer](#), [\[2011.00416\] Deep Learning for Text Style Transfer: A Survey](#)
12. Authorship attribution / stylometry
- Ideas: automatically identify author of literary texts (aim for similar authors) / classify characters in a play; automatically identify authors of social media posts (e.g. authorship attribution/verification on tweets); automatically classify authors of scientific papers; identify authors of song lyrics; + analyze which feature characterize the authors (explainability of the machine learning model)
  - Papers & authors: [PAN](#), [Liviu Dinu](#), [Marius Popescu](#)
13. Plagiarism detection
- Ideas: implement a plagiarism detection system based on a collection of scientific articles (identify duplicate content and source); solve a shared task on plagiarism detection (PAN); plagiarism detection on political speeches
  - Papers & authors: [PAN](#), [Liviu Dinu](#), [Marius Popescu](#)
14. Author profiling (detecting the age/gender/personality/native language of an author)
- Ideas: profile users on social media (automatically predict demographics like age/gender/location);
  - Papers & authors: PAN, [Sulea& Dichiu](#), [Sulea&Zampieri](#), [Nisioi](#)
15. Hate speech, offensive language identification, misogyny / stereotype detection
- Ideas: annotate a corpus of Romanian tweets with hate speech labels; solve a shared task on hate speech/aggressive speech detection;
  - Sexism, Racism, Homophobia <https://arxiv.org/pdf/2106.15896.pdf>
  - Papers & authors: [Marcos Zampieri](#), [Paolo Rosso](#)
16. Computational humor, sarcasm & irony detection
- Ideas: implement an irony/sarcasm detector on social media, predict the #irony/#sarcasm hashtag; automatically identify sarcasm in movie or product reviews/news; measure correlation between sarcasm and optimism/pessimism/emotions expressed or author personality/profile on social media texts
  - Papers & authors: [A multidimensional approach for detecting irony in Twitter](#) , [Cristian Danescu Mizil](#), Carlo Straparava, [Paolo Rosso](#)
17. Metaphor and figurative language detection

- Ideas: automatically identify words used metaphorically in poems/song lyrics/social media; automatically identify meaning of a metaphor (“translate” the metaphor);
  - Papers: [Brighter than Gold: Figurative Language in User Generated Comparisons](#), [From humor recognition to irony detection: The figurative language of social media](#), [A Computational Exploration of Exaggeration](#), [Impact Analysis of Emotion in Figurative Language](#)
18. Diachronic and historical linguistics: word formation, cognates identification, proto-word re-construction, borrowing, language similarity, etc
- Ideas: collect lexicon of words and etymologies in a low resource language (Eastern European languages?); compare different measures of language/dialect similarity based on common vocabulary/similar syntax/similar phonetics;
  - Papers & authors: [Liviú Dinu](#), [Alina Maria Ciobanu](#)
19. Semantic change – tracking the change in meanings of words
- Ideas: identify semantic change in certain subset of words e.g. sentiment/emotion words (optional: for Romanian), in business terminology; compare metaphorical senses of words across languages; identify changes in slang terms and appearance of new senses from social media data (e.g. “lit”); solve shared task on semantic change (SemEval 2020)
  - Papers: [Towards Computational Lexical Semantic Change Detection](#) (LChange Workshop - check proceedings); [Computational approaches to semantic change](#) book
20. Temporal text classification, dating of texts
- Ideas: predict period when text was written based on different features; identify which features are most useful for dating: news texts/scientific texts/
  - Papers: [Temporal classification for historical Romanian texts](#)
21. Law and NLP-AI
- Ideas: predict outcome of court cases; automatically parse contracts
  - Papers: [\[1710.09306\] Exploring the Use of Text Classification in the Legal Domain](#)
22. NLP and ethics, biases in datasets and algorithms (explainability/interpretability)
- Ideas: identify biases against immigrants in news texts using word embeddings (better: contextual embeddings) - biased sentiment/emotion; track changes in biases wrt certain minority over time; evaluate bias in Romanian embeddings (towards minorities in Romania?)
  - Papers: [Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings](#),
23. Language and text generation; BERT, GPT-3 and other Transformer based language models
- Ideas: build a language generator for Romanian based on pre-trained transformers (e.g. multilingual BERT); generate (fake?) news/scientific articles/food recipes
  - Examples: <https://app.inferkit.com/demo>, <https://transformer.huggingface.co/>, <https://6b.eleuther.ai/>

- Papers: [\[1810.04805\] BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#), <https://aclanthology.org/2020.coling-main.581.pdf>
  - Resources: [https://huggingface.co/transformers/pretrained\\_models.html](https://huggingface.co/transformers/pretrained_models.html), <https://github.com/kingoflolz/mesh-transformer-jax/#gpt-j-6b>
24. Social Media text processing (fb, ig, twitter, reddit)
- Ideas: any task applied to social media text: hate speech, word embeddings (identify new terms), authorship, author profiling, opinion mining; tracking emotions expressed on social media (bonus: do it for Romanian - more rare); collect and annotate a corpus of social media data (ideally for a new language/task); meme analysis - multimodal text+visual (papers: [Analysis of Facebook Meme Groups Used During the 2016 US Presidential Election](#), [Findings of the WOAHS 5 Shared Task on Fine Grained Hateful Memes Detection](#), [Competition Memotion](#)) (pretrained image analysis: VGG...)
25. Deep learning in NLP
- Ideas: any NN architecture for any NLP task, explain the math
26. Machine learning in NLP (supervised, weakly supervised, zero shot learning, few shot learning; probabilistic models, variational inference)
- Ideas: non-DL machine learning for an NLP task, compare models
27. Transfer learning – multi-stage machine learning where knowledge from one dataset/domain/task is leveraged to help with another
- Papers & authors: Sebastian Ruder
28. (Neural) Machine Translation
- Ideas: implement a machine translation model
  - Papers & tools: any Google Translate paper, Interlingua, WMT
29. Word sense disambiguation
- Ideas: implement a word sense disambiguation model (bonus: for Romanian/a low-resource language); evaluate contextual embeddings (transformer-based) for disambiguation, compare with other models
  - Papers & authors: Florentina Hristea
30. Morphological re-inflection; Inflection generation
- Ideas: reproduce models on inflection generation from existing papers on Romanian, apply for new parts of speech / use different architectures
  - Papers & authors: Maria Sulea, SIGMORPHON
31. Collocation detection; multi word expressions, phrase identification
- Ideas: implement/compare collocation detection models; compare meanings of phrases in different languages; embeddings for collocations/phrases
  - Papers & authors: Mikolov, Katja Lipshikova
32. Anaphora resolution, coreference resolution
- Ideas: implement model that identifies anaphora; implement a tool to visualize coreference in texts
  - Papers: [Global Inference for Bridging Anaphora Resolution - Yufang Hou<sup>1</sup>, Katja Markert<sup>2</sup>, Michael Strube<sup>1</sup>](#), [An Algorithm for Pronominal Anaphora Resolution](#),
33. Sentiment Analysis; optimism-pessimism identification

- Ideas: implement +evaluate an algorithm for sentiment analysis - predict sentiment expressed in tweet / news / review (optional: use Romanian data); optional: identify specifically sentiment for individual aspects of the object (aspect-based sentiment analysis)
- Papers & authors: [Exploiting BERT for end-to-end aspect-based sentiment analysis](#) , Rada Mihalcea

#### 34. Opinion Mining

- Ideas: implement an algorithm for understanding the opinion on a given product / service / public person from online reviews / social media (e.g. on iPhone, PNL, vaccine, ...) (optional: use Romanian data); preliminary: annotate a social media dataset with sentiment scores / scrape reviews and annotate based on stars / use existing dataset; optional: identify specifically which aspects of the product are being referred to - see aspect-based sentiment analysis (e.g. camera is good, battery is bad)

#### 35. Emotion analysis – detect the emotions in a text

- Ideas: automatically extract emotion scores for individual emotions (see Plutchik's wheel of emotions) for news / blogs / social media texts; optional: use Romanian data; optional: analyze emotions with respect to a given label for a separate task, or use as feature for separate task: e.g. emotions in hate speech datasets, emotions in optimism/pessimism data, emotions in therapy sessions; tracking emotions on social media over time; annotate a dataset with emotion scores and learn to predict them (preferably for a low-resource language/new domain)
- Resources: <https://saifmohammad.com/WebPages/lexicons.html>

#### 36. NLP for clinical/ medical data

- Ideas: NER model for clinical data; information retrieval in medical texts; build embeddings for medical terminology
- Papers: [NER for Medical Entities in Twitter using Sequence to Sequence Neural Networks](#), [Adaptive Generation of Structured Medical Report Using NER Regarding Deep Learning](#), Proceedings of [https://aclweb.org/aclwiki/BioNLP\\_Workshop](https://aclweb.org/aclwiki/BioNLP_Workshop) , pubmed

#### 37. <https://naacl2018.wordpress.com/2018/01/14/test-of-time-paper-nominations-or-classic-computational-linguistics-papers/> (discutia unui articol din aceasta lista)

#### 38. Recent research topics in NLP (articole recente relevante din Computational Linguistics, ACL, COLING, EMNLP, NAACL, EACL, PNAS, etc) see best papers proposals in the last 10 years (Disponibile on-line la [https://aclweb.org/aclwiki/Best\\_paper\\_awards](https://aclweb.org/aclwiki/Best_paper_awards) )

#### 39. LREC 2020 papers for re-experimentation

(<https://lrec2020.lrec-conf.org/en/reprolang2020/selected-tasks/>).

#### 40. NLP applications

- Ideas: resume analysis, automatic question tagging on StackOverflow/Quora etc, spam classification, automatic essay grading, bot detection, recommender system for products, movies etc

#### 41. NLP & Art



- Ideas: lyrics generation (constrained to rhyme?); classification of literary texts/poems/song lyrics; generate text in the style of Shakespeare...
  - Papers: <https://rootroo.com/en/hucmac/>, [Creative GANs for generating poems, lyrics, and metaphors](#), [Weird AI Yankovic: Generating Parody Lyrics](#)
42. NLP for literary texts / digital humanities
- Ideas: profiling literary characters, character networks, detecting events, profiling authors based on literary texts, OCR for historical texts; build a text processing tool to assist linguists/historians/etc...
  - Resources: <https://www.gutenberg.org/>
  - Papers: <https://sighum.wordpress.com/> (look at proceedings), <https://text2story22.inesctec.pt/>
43. Search engine
- Ideas: implement a search engine / information retrieval system on a corpus of data; implement application to allow users to perform searches
44. Text to Speech, Speech to Text
- Ideas: speech2text system for a chatbot
45. Mental health, depression detection, etc
- Ideas: solve an eRisk challenge/a CLPsych challenge (free datasets, you need to request the datasets from the organizers); collect and annotate a corpus on depression/another mental illness for a low-resource language
  - Papers: <https://erisk.irlab.org/>, <https://clpsych.org/> (see proceedings)
  - Datasets:
    - i. [kharrigian/mental-health-datasets: An evolving list of electronic media data sets used to model mental-health status.](#)
    - ii. Anorexia. Data from [Early risk prediction on the Internet | CLEF 2019 workshop](#)
    - iii. Self-harm. Data from [CLEF eRisk: Early risk prediction on the Internet | CLEF 2021 workshop](#)
    - iv. PTSD [Measuring Post Traumatic Stress Disorder in Twitter - Glen Coppersmith Craig Harman Mark Dredze](#)
    - v. Suicide Ideation (hard to get access datasets due to ethical concerns)
    - vi. Bipolar disorder [Not Just Depressed: Bipolar Disorder Prediction on Reddit](#)
    - vii. Stress [Dreaddit: A Reddit Dataset for Stress Analysis in Social Media](#)
    - viii. Multiple Mental Health Conditions Classification [SMHD: A Large-Scale Resource for Exploring Online Language Usage for Multiple Mental Health Conditions](#)
46. Conversational agent / chatbot
- Ideas: implement a chit-chat bot, customer support bot, Alexa-style bot, robot therapist... , optional: integrate with Google API for complete app including speech module
  - Papers: [An Overview of Chatbot Technology](#), [Chatbot for university related FAQs](#), [Conversational agents in healthcare: a systematic review | Journal of the American Medical Informatics Association | Oxford Academic](#),

[https://www.academia.edu/download/61052907/reportmedical\\_chatbot20191028-44031-hq2g2a.pdf](https://www.academia.edu/download/61052907/reportmedical_chatbot20191028-44031-hq2g2a.pdf)

- Examples: Amazon Alexa, Google Assistant, [Eliza, a chatbot therapist](https://www.talktopoppy.com/), <https://www.talktopoppy.com/>
47. Topic modeling – extract topics discussed in a text (classical LDA / neural topic modeling)
- Ideas: identify and track topics over time in news/scientific texts/social media; implement a topic model from scratch; dynamic topic modelling
  - Papers: [Dynamic Topic Models](#), [A Novel Approach of Neural Topic Modelling for Document Clustering](#), [Discovering Discrete Latent Topics with Neural Variational Inference](#), [Studying the Evolution of Scientific Topics and their Relationships](#)
48. Image captioning (automatically generate a description of an image – involves both NLP and computer vision)
- Ideas: implement+evaluate a model that generates image captions, ideally in sentence format (focus on text generation part)
  - Papers: [Convolutional Image Captioning](#), [A Comprehensive Survey of Deep Learning for Image Captioning](#)
49. Language identification
- Ideas: implement model that identifies language/dialect of given text, code switching detection (“romgleza” etc)
  - Papers: [Automatic Language Identification in Texts: A Survey](#) , [VarDial 2019 - Evaluation Campaign](#) , [Proceedings of the 8th VarDial Workshop on NLP for Similar Languages, Varieties and Dialects](#), [Code-switching detection using multilingual DNNS | IEEE Conference Publication](#), [Recurrent-neural-network for language detection on twitter code-switching corpus](#)
50. Question answering
- Ideas: implement question answering model
  - Papers: [QuAC: Question Answering in Context](#) , [The Question Answering Systems: A Survey](#).
51. Stance detection, hyperpartisanship etc
- Ideas: predict political views, stance on vaccine / social matters / etc; solve shared task on stance detection, bias in news
  - Papers: [Multi-Task Stance Detection with Sentiment and Stance Lexicons](#) , <http://snap.stanford.edu/quotus/#about>
  - Datasets: [An Interactive Visualization of the SemEval-2016 Stance Dataset](#) , [Stance Detection](#)
52. Explainable AI for NLP
- Ideas: implement explainability methods for a NLP machine learning model (e.g. attention weights analysis, LIME, adversarial examples, gradients analysis...)
  - Papers: [\[2009.13295\] A Diagnostic Study of Explainability Techniques for Text Classification](#)
  - Resources: <https://github.com/marcotcr/lime>, <https://tf-explain.readthedocs.io/en/latest/>
53. Other NLP & CL topics (send an email for approval)



## Requirements & guidelines:

Projects should consist of 3 parts:

- paper / technical report
- implementation (code)
- slide presentation
- + a short document explaining the contribution of each student in the team

A project can be focused either on:

- the paper (**survey**) – detailed presentation of existing methods; in this case the paper should be focused on describing the state of the art, comparing existing methods + should contain a proof of concept implementation of a solution to the chosen problem,
- or the **implementation** (in this case the paper will be focused on the methodology and describing technical and experimental details):
  - implementing an end-to-end application to solve the problem
  - implementing a solution described in a paper (projects 37, 38, 39)
  - implementing a novel solution to a problem

All papers/technical reports will follow the classical structure of a research article (approx 4 pages):

- short summary (abstract)
- analysis of main idea
- related work: state of the art (SOTA) where it exists, short history, recent and/or related results
- In case you're presenting a survey: explain main methodologies and selection process (i.e. you are surveying either chronologically, or in order of SOTA achievements), discuss advantages and disadvantages to the methods used and introduced
- in the case of presenting specific applications: describe the method, compare it with other results in the field
- conclusions and future work, directions for further improvement
- references

Teams of 2-3, max 4 people. Any topic can be chosen by max 4 teams.

Add your name on the google sheets document next to the chosen project topic.

Additional details on useful resources here:

<https://github.com/anana/nlp-projects/blob/main/README.md>